



基于k8s的Flink SQL大 规模数据相似度计算

汇报人：22121630 汪江豪
22120721 阮金桐

C O N T E N T

目 录

01

Flink多节点部署方式

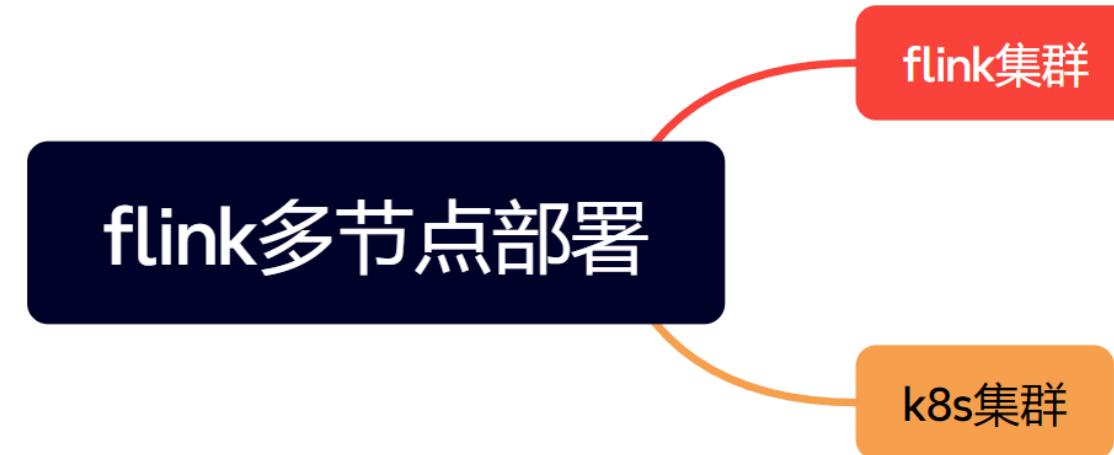
02

解法思路

03

问题处理

1. Flink多节点部署方式



传统集群方式

设置ssh无密码互联

修改flink/conf/flink-conf.yaml文件:

obmanager.rpc.address: master

taskmanager.numberOfTaskSlots: 2

修改conf目录下master文件为
master:8081

修改conf目录下worker文件为:

slave1

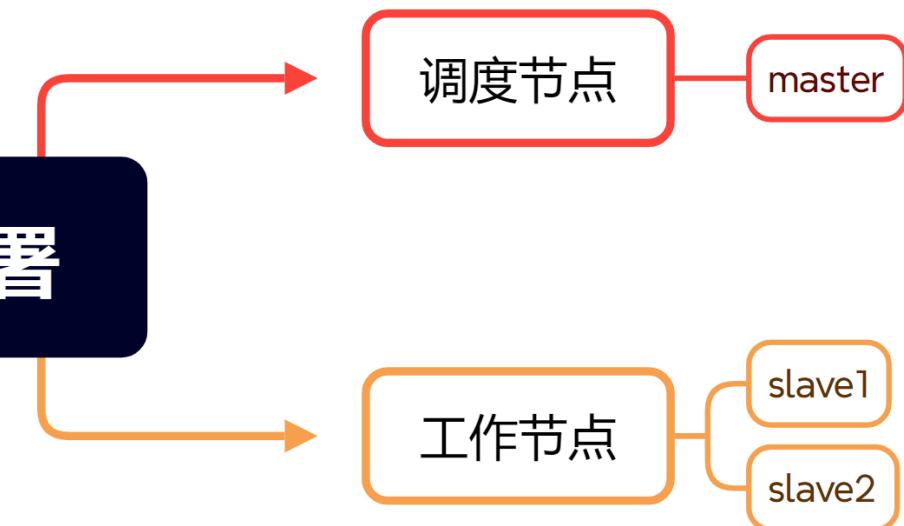
slave2

slave1,slave2节点上设置同理

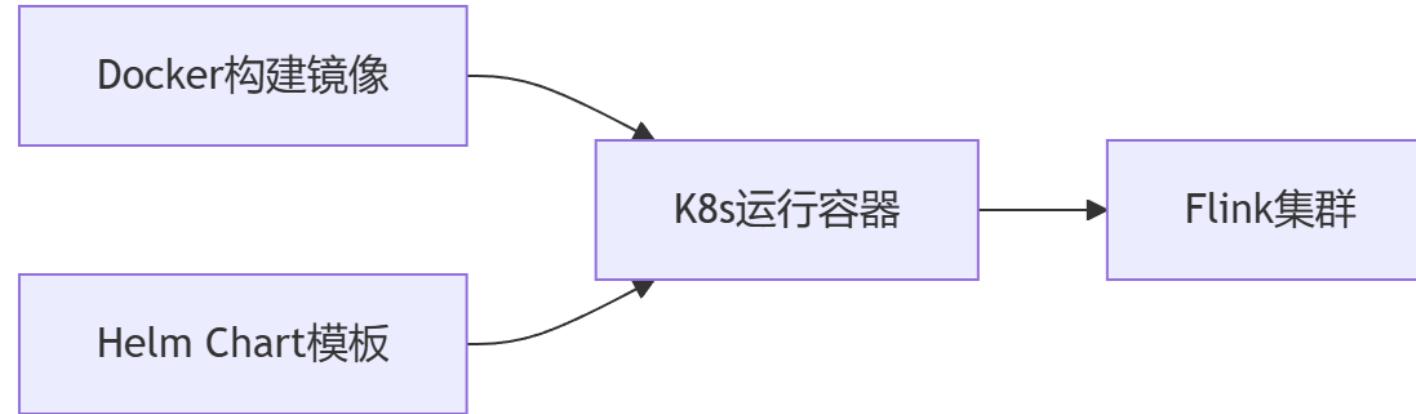
额外修改flink-conf.yaml文件中的
taskmanager.host:slave1/slave2
为自己主机名

在master节点运行start-cluster.sh, 提交
jar包

传统集群部署



依赖项



K8s是容器编排平台，用于自动化部署、扩展和管理容器化应用

Flink中的jobmanager和taskmanager作为pod运行在k8s集群中

Docker是容器化工具，可将应用及其依赖打包成轻量级、可移植容器镜像

Helm是k8s的包管理工具，类似于linux的apt

K8s集群方式部署Flink



启动k8s集群，节点加入集群

下载证书管理器

```
kubectl create -f https://github.com/jetstack/cert-manager/releases/download/v1.8.2/cert-manager.yaml
```

添加软件源

```
helm repo add flink-operator-repo https://downloads.apache.org/flink/flink-kubernetes-operator-1.11.0/
```

安装flink的k8s管理工具

```
helm install flink-kubernetes-operator flink-operator-repo/flink-kubernetes-operator
```

安装后operator作为pod保持运行，看到running说明安装成功

```
kubectl get pods
```

随后可通过yaml文件提交flink作业

```
Kubectl create -f basic.yaml
```

若要查看flink的webui，可进行端口转发，并允许外来地址访问

```
kubectl port-forward svc/largejob-rest 8081 --address 0.0.0.0
```

basic.yaml文件



```
1 apiVersion: flink.apache.org/v1beta1
2 kind: FlinkDeployment
3 metadata:
4   namespace: default
5   name: largejob
6 spec:
7   image: flink:1.16
8   flinkVersion: v1_16
9   flinkConfiguration:
10    taskmanager.numberOfTaskSlots: "12"
11   serviceAccount: flink
12   jobManager:
13     resource:
14       memory: "4096m"
15       cpu: 2
16   podTemplate:
17     spec:
18       # nodeSelector:
19         # flink-role: jobmanager
20     tolerations:
21       - key: "node-role.kubernetes.io/master"
22         operator: "Exists"
23         effect: "NoSchedule"
24     containers:
25       - name: flink-main-container
26         volumeMounts:
27           - name: my-jar
28             mountPath: /opt/flink/largejob
29     volumes:
30       - name: my-jar
31         hostPath:
32           path: /home/ubuntu/tableapp/upload
33           type: Directory
34           gType: Directory
35           path: \home\ubuntu\tableapp\upload
36           type: Directory
```

```
4 taskManager:
5   resource:
6     memory: "51200m"
7     cpu: 12 # 每个taskmanager分配12个CPU
8   replicas: 3 # 启动2个taskmanager pod, 调度到2个节点
9   podTemplate:
10     spec:
11       nodeSelector:
12         flink-role: taskmanager
13       tolerations:
14         - key: "node-role.kubernetes.io/master"
15           operator: "Exists"
16           effect: "NoSchedule"
17       containers:
18         - name: flink-main-container
19           volumeMounts:
20             - name: my-jar
21               mountPath: /opt/flink/largejob
22       volumes:
23         - name: my-jar
24           hostPath:
25             path: /home/ubuntu/tableapp/upload # 本地JAR所在目录
26             type: Directory
27   job:
28     jarURI: local:///opt/flink/largejob/inputfromfiledemo-1.0.jar
29     entryClass: cn.edu.shu.large_result # 填写作业的主类
30     parallelism: 36
```

将本地数据集和jar包放在所有节点同一目录下，并修改权限为777
挂载到flink镜像中的某个目录

这里本地数据集和jar包位于/home/ubuntu/tableapp/upload

挂载目录为flink镜像中的/opt/flink/largejob

2. 解法思路

解决办法

方法	配置	实际使用资源	总计耗时
Flink集群	3台4核8G 2个TM节点	2*4核8G	34m50s
K8s集群	3台16核64G	3*12核50G	8m50s

解法思路

1	116420525701620762252	106501936839371135489
2	104979585499014359063	103039370087174812731
3	101130571432010257177	111183544898345861357
4	116331515612347682756	114390577443742844396
5	110656253137238747097	102565924973578852934
6	109800388691366698136	102545157386069758716
7	106752695486123789059	115120856388820348743
8	117594348100980996964	100517144772812557903
9	109405659400238396060	105452294703789324242
10	102857824121129353719	104147032621576433597

约7万个

约10万个



原始数据特点：
6832726行、21位数字字符串、无重复行

存储类型：
INT：10位×
BIGINT：19位×
STRING、DECIMAL性能开销大×

解决办法——转换数据集：

排序——两列数据分别映射到0-7万，0-10万

保存映射字典，供后续转换使用

使用INT存储

映射字典：

第一列

large_referrer_dict.csv	
68196	118441596793086402969,68194
68197	118441827513137270649,68195
68198	118441866522954267546,68196
68199	118443887616207357079,68197
68200	118443964499336832769,68198
68201	118444229992664950227,68199
68202	118444248646004441122,68200
68203	118444997653815563102,68201
68204	118445277475148270304,68202
68205	118445438046067769625,68203
68206	118446147598193798543,68204
68207	118446153145043212994,68205
68208	118446153145043212994,68205

原始数据集：286M

第二列

large_referree_dict.csv	
101210	118444005060455277166,101208
101211	118444229992664950227,101209
101212	118444401187415650231,101210
101213	118444576709377182815,101211
101214	118444881207894739198,101212
101215	118444997653815563102,101213
101216	118445277475148270304,101214
101217	118445438046067769625,101215
101218	118446147598193798543,101216
101219	118446153145043212994,101217
101220	118446297469672800612,101218
101221	118446413811161034753,101219
101222	118446413811161034753,101219

转换后数据集83M



large_relation.csv	
1	60777,35921
2	18389,16589
3	4084,61626
4	60452,79054
5	39424,13946
6	36298,13838
7	25039,83137
8	65162,2781
9	34885,30107
10	10436,22831
11	50032,38691
12	20105,43662
13	56302,55433
14	1755,60672
15	49336,4561
16	37054,55332
17	406,62598
18	6139,88632
19	54334,86173
20	15570,5331
21	36673,34743

JAVA代码



在flink1.12之后，批处理模式通常比流处理模式在处理静态大数据集时效率更高：

批处理模式会自动做全局优化如(全局排序、资源复用等)。

流处理模式主要为低延迟，适合实时场景。

在传统集群方式中，
java代码设置并行数以提高效率

K8s集群方式，
java代码不用设置并行数，
在yaml文件中指定

JAVA代码续

对于传统flink集群方式，path为本地数据集路径。

对于k8s集群方式，path为flink镜像中的挂载目录

```
// 结果输出到文件系统
tableEnv.executeSql("CREATE TABLE SinkTable (
    + "web1 INT, "
    + "web2 INT, "
    + "similarity DOUBLE"
    + ") WITH (
    + 'connector' = 'filesystem',
    + 'path' = 'file:///opt/flink/largejob/large_result',
    + 'format' = 'csv',
    + 'csv.field-delimiter' = ',',
    + 'sink.rolling-policy.file-size' = '256MB',
    + 'sink.rolling-policy.rollover-interval' = '30 min'
    + ")");
+ ..)
+ ..,
+ ..,
+ ..,
```

```
// 创建临时视图 ref_count
tableEnv.executeSql("CREATE TEMPORARY VIEW tmp_ref_count AS "
    + "SELECT referer, COUNT(DISTINCT referree) AS web_count "
    + "FROM large_relation "
    + "GROUP BY referer");

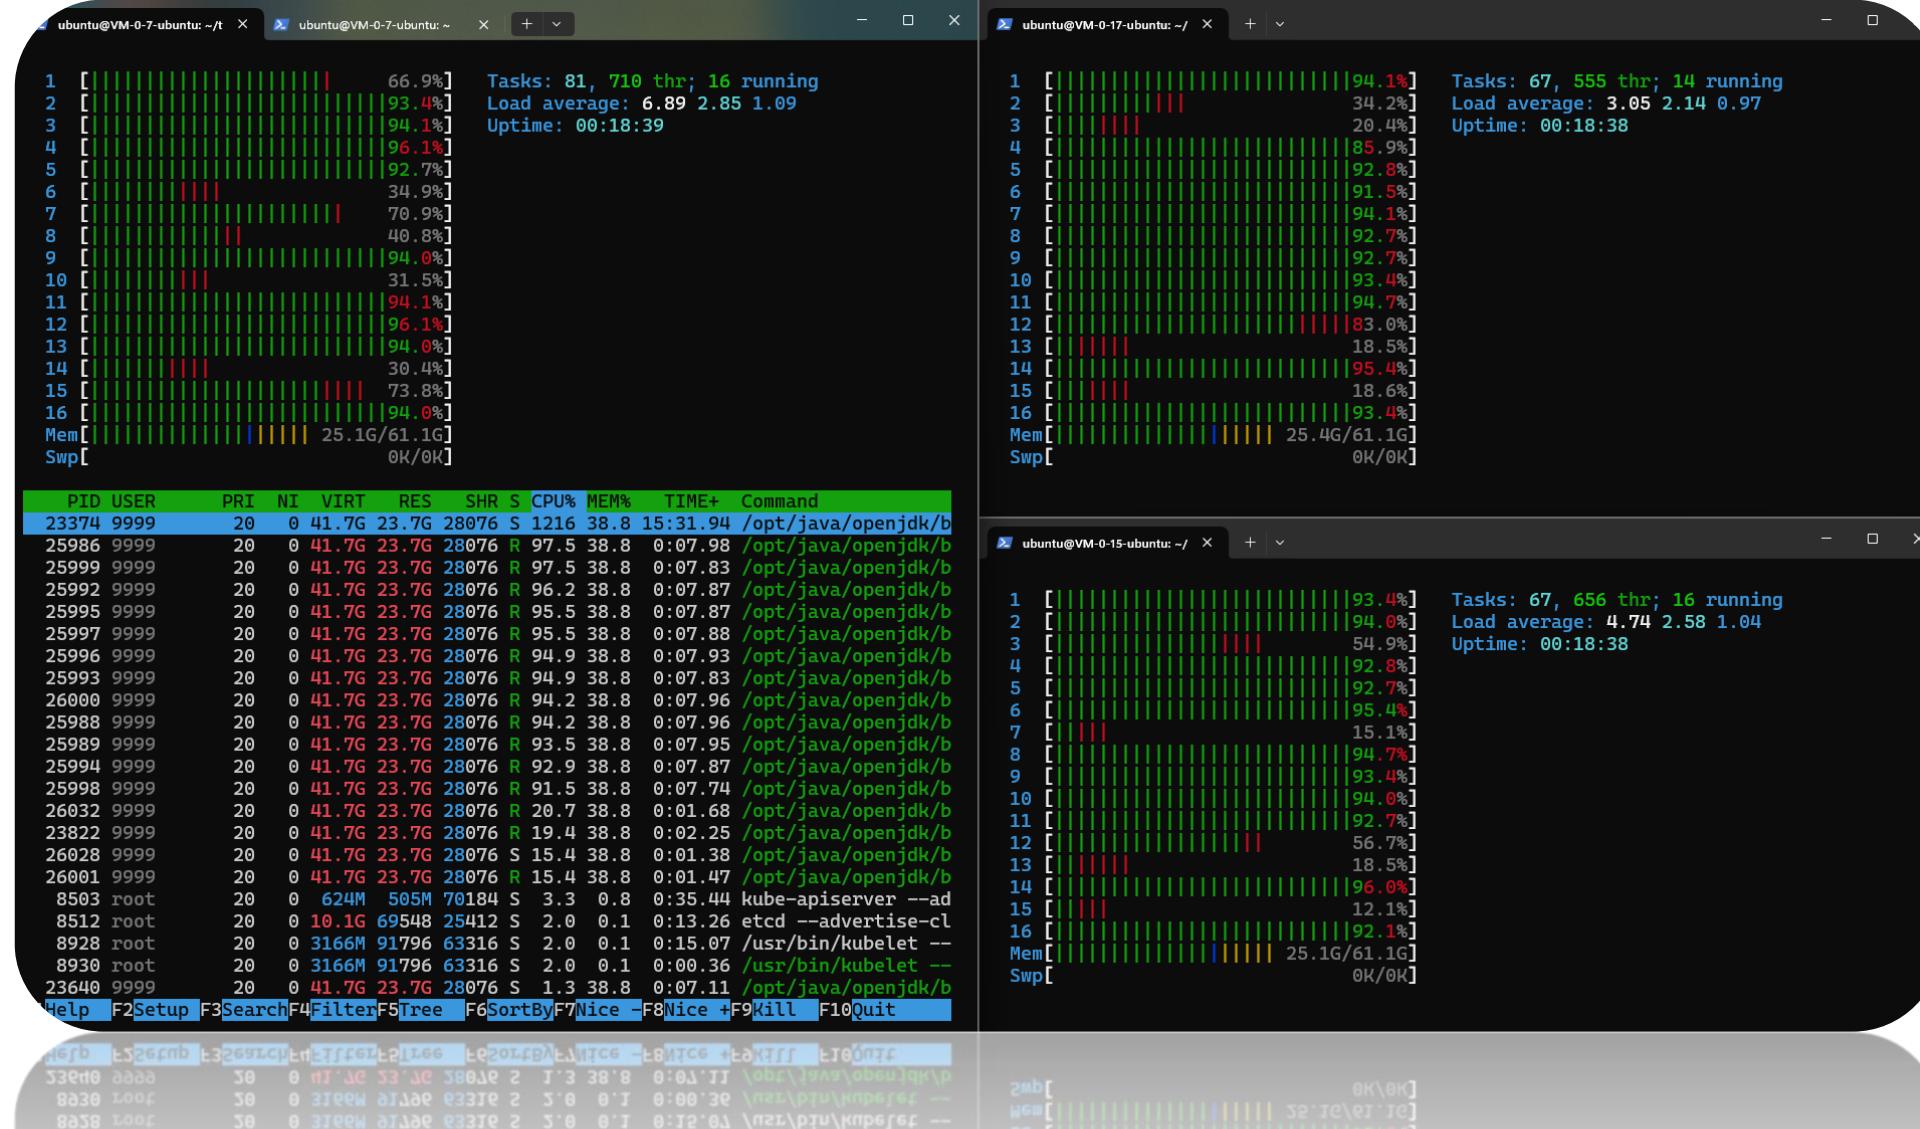
// 创建临时视图 common
tableEnv.executeSql("CREATE TEMPORARY VIEW tmp_common AS "
    + "SELECT a.referrer AS web1, b.referrer AS web2, COUNT(*) AS com_cnt "
    + "FROM large_relation a "
    + "JOIN large_relation b ON a.referee = b.referee "
    + "WHERE a.referrer < b.referrer "
    + "GROUP BY a.referrer, b.referrer");

// 计算相似度并插入结果表
String insertSQL = "INSERT INTO SinkTable "
    + "SELECT "
    + "    common.web1 AS web1, "
    + "    common.web2 AS web2, "
    + "    CASE "
    + "        WHEN r1.web_count + r2.web_count - common.com_cnt > 0 "
    + "        THEN common.com_cnt * 1.0 / (r1.web_count + r2.web_count - common.com_cnt) "
    + "        ELSE 0 "
    + "    END AS similarity "
    + "FROM tmp_common common "
    + "JOIN tmp_ref_count r1 ON common.web1 = r1.referrer "
    + "JOIN tmp_ref_count r2 ON common.web2 = r2.referrer";
tableEnv.executeSql(insertSQL);
```

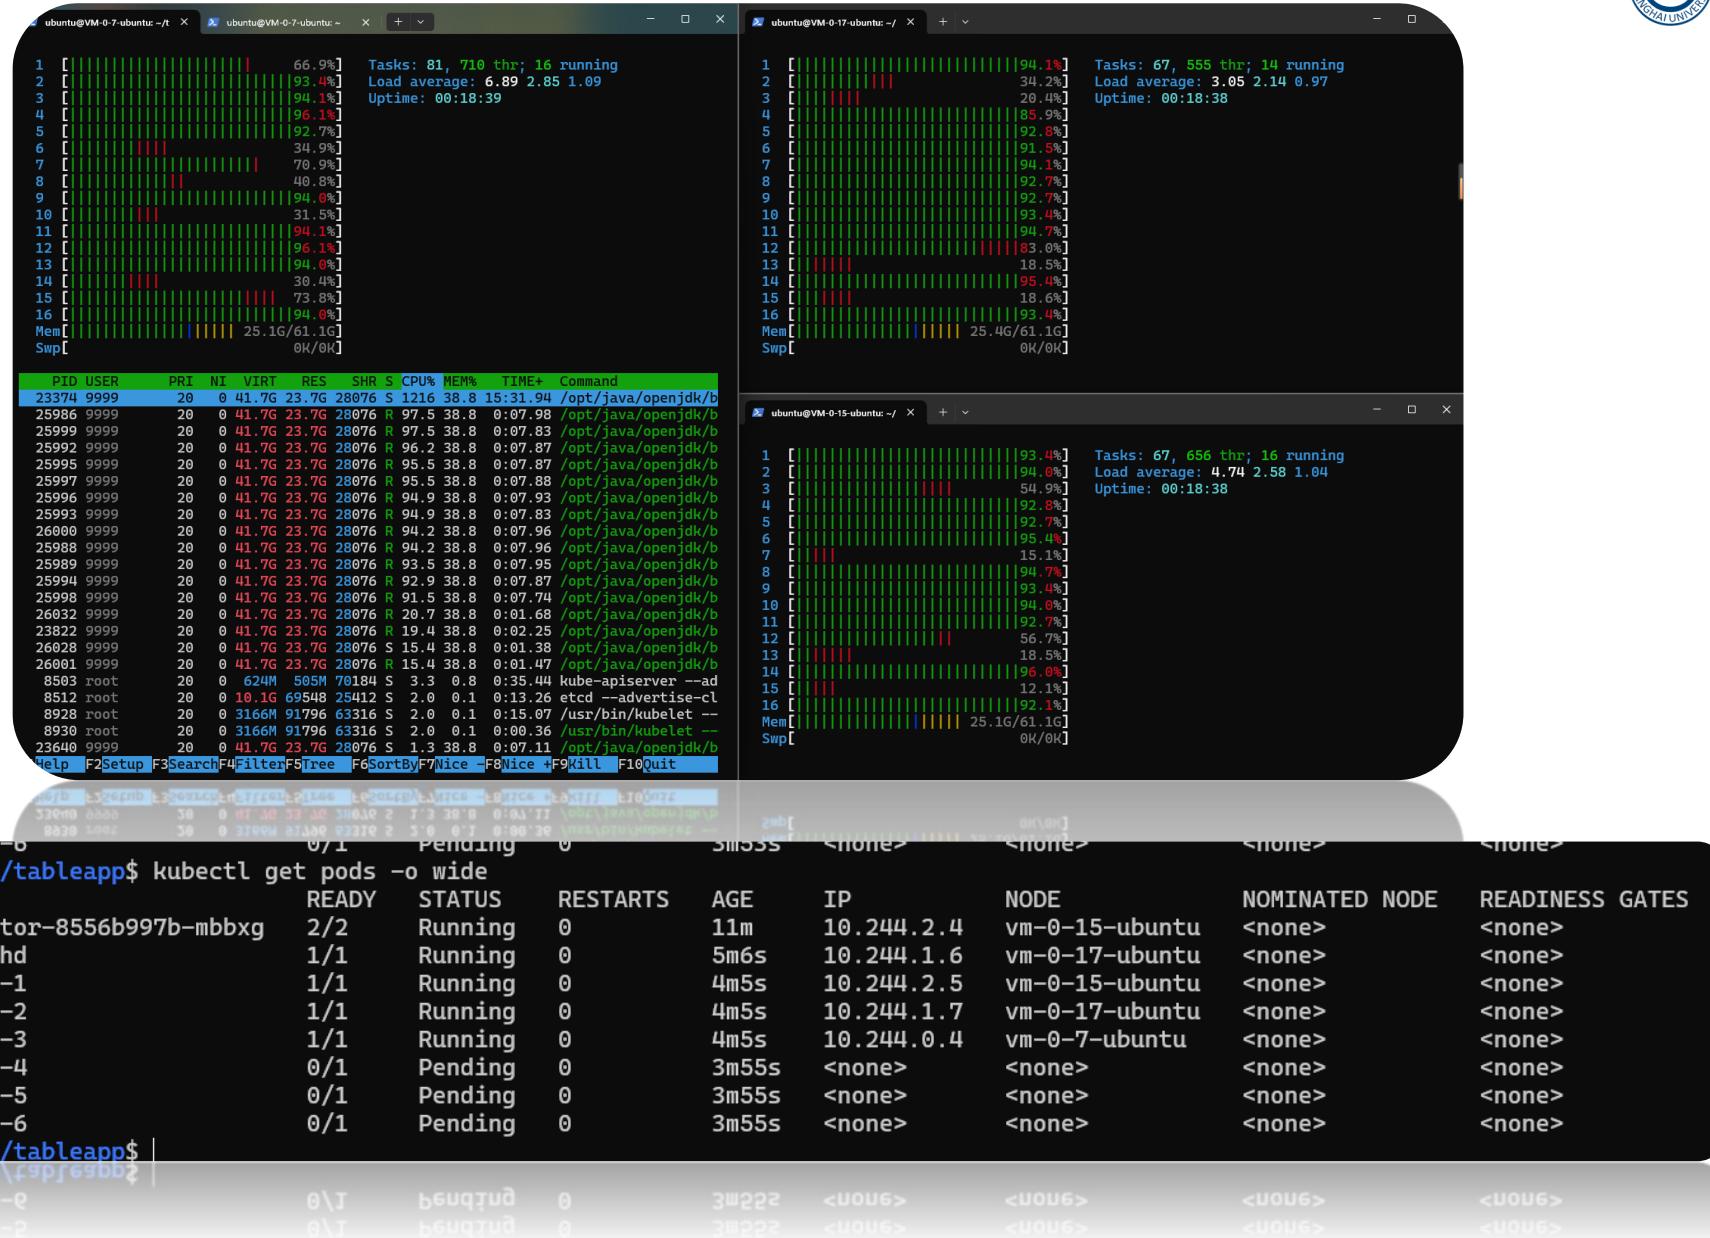
在basic.yaml文件中，将数据集和jar的目录挂载好后，通过命令：

kubectl create -f basic.yaml创建任务，会自动创建镜像，执行查询

运行截图



运行截图



运行截图

Apache Flink Dashboard

Version: 1.16.3 Commit: 5df8156 @ 2023-11-10T11:00:33+01:00 Message:

largejob

Job ID	79c6bb41c1ac7db0acaffe72d65eae	Job State	FINISHED 7	Actions	Job Manager Log
Start Time	2025-05-19 15:34:53	End Time	2025-05-19 15:40:29	Duration	5m 35s

Overview Exceptions Timeline Configuration

The job graph shows the flow of data between various tasks. The tasks are represented as boxes with their names, parallelism, and resource usage. Edges between tasks indicate the direction of data flow, labeled with 'FORWARD' or 'BACKWARD'. The tasks include Source, HashAggregate, HashJoin, Calc, and Sink.

Name	Status	Bytes Received	Records Received	Bytes Sent	Records Sent	Parallelism	Start Time	Duration	End Time	Tasks
Source: large_relation[1] -> HashAggregate[9]	FINISHED	0 B	0	483 MB	6,836,726	36	2025-05-19 15:34:57	6s	2025-05-19 15:35:04	36
HashAggregate[11] -> HashAggregate[12]	FINISHED	215 MB	6,836,726	38.7 MB	1,229,301	36	2025-05-19 15:35:04	1s	2025-05-19 15:35:05	36
HashAggregate[14]	FINISHED	38.7 MB	1,229,301	2.15 MB	68,206	36	2025-05-19 15:35:53	1s	2025-05-19 15:35:55	36
HashJoin[3] -> Calc[4] -> HashAggregate[5]	FINISHED	535 MB	13,673,452	97.2 GB	2,546,776,700	36	2025-05-19 15:35:05	3m 2s	2025-05-19 15:38:08	36
HashAggregate[7]	FINISHED	97.2 GB	2,546,776,700	14.6 GB	383,157,187	36	2025-05-19 15:38:08	1m 7s	2025-05-19 15:39:16	36
HashJoin[15] -> Calc[16]	FINISHED	14.6 GB	383,225,393	17.5 GB	383,157,187	36	2025-05-19 15:39:16	8s	2025-05-19 15:39:24	36
HashJoin[18] -> Calc[19] -> Sink: Filesystem	FINISHED	17.5 GB	383,225,393	0 B	0	36	2025-05-19 15:39:24	1m 5s	2025-05-19 15:40:29	36

结果处理——生成36个子文件

The image shows three terminal windows side-by-side, each displaying the output of the `ls` command in a directory named `large_result/.staging_1747640093663`. The leftmost window (VM-0-7) lists files named task-12 through task-35. The middle window (VM-0-17) lists files named part-e72668a2-1e93-413e-9898-6578f6880474-task-0-file-0 through part-e72668a2-1e93-413e-9898-6578f6880474-task-9-file-0. The rightmost window (VM-0-15) lists files named task-13 through task-34.

```
ubuntu@VM-0-7-ubuntu:~/tableapp/upload/large_result/.staging_1747640093663$ ls
task-12  task-17  task-21  task-25  task-29  task-33
task-14  task-20  task-24  task-26  task-31  task-35
ubuntu@VM-0-7-ubuntu:~/tableapp/upload/large_result/.staging_1747640093663$
```

```
ubuntu@VM-0-17-ubuntu:~/tableapp/upload/large_result/.staging_1747640093663$ ls
part-e72668a2-1e93-413e-9898-6578f6880474-task-0-file-0
part-e72668a2-1e93-413e-9898-6578f6880474-task-10-file-0
part-e72668a2-1e93-413e-9898-6578f6880474-task-11-file-0
part-e72668a2-1e93-413e-9898-6578f6880474-task-1-file-0
part-e72668a2-1e93-413e-9898-6578f6880474-task-2-file-0
part-e72668a2-1e93-413e-9898-6578f6880474-task-3-file-0
part-e72668a2-1e93-413e-9898-6578f6880474-task-4-file-0
part-e72668a2-1e93-413e-9898-6578f6880474-task-5-file-0
part-e72668a2-1e93-413e-9898-6578f6880474-task-6-file-0
part-e72668a2-1e93-413e-9898-6578f6880474-task-7-file-0
part-e72668a2-1e93-413e-9898-6578f6880474-task-8-file-0
part-e72668a2-1e93-413e-9898-6578f6880474-task-9-file-0
ubuntu@VM-0-17-ubuntu:~/tableapp/upload/large_result/.staging_1747640093663$
```

```
ubuntu@VM-0-15-ubuntu:~/tableapp/upload/large_result/.staging_1747640093663$ ls
task-13  task-16  task-19  task-23  task-28  task-32
task-15  task-18  task-22  task-27  task-30  task-34
ubuntu@VM-0-15-ubuntu:~/tableapp/upload/large_result/.staging_1747640093663$
```

结果处理——显示结果文件信息

```
ubuntu@VM-0-7-ubuntu:~/all_tasks$ ls
task0  task10  task12  task14  task16  task18  task2  task21  task23  task25  task27  task29  task30  task32  task34  task4  task6  task8
task1  task11  task13  task15  task17  task19  task20  task22  task24  task26  task28  task3  task31  task33  task35  task5  task7  task9
```

```
ubuntu@VM-0-7-ubuntu:~/all_tasks$ wc -l task
383157187 task
```

```
ubuntu@VM-0-7-ubuntu:~/all_tasks$ ls -l
total 11637768
```

```
-rw-rw-r-- 1 ubuntu ubuntu 11917070025 May 19 16:12 task
```

```
ubuntu@VM-0-7-ubuntu:~/all_tasks$ head -n 10 task
```

```
25556,49046,0.02527075812274368
2922,35245,0.01211556383970177
60482,64263,0.06116207951070336
49686,60622,0.02987551867219917
4769,60665,0.00836820083682008
4769,36941,0.016666666666666667
4769,57212,0.0101010101010101
57749,67620,0.0303030303030303
2072,5687,0.03486238532110092
4769,65581,0.0108695652173913
```

结果处理——Python脚本进行反射

```
def retransform_dict():
    start_time = time.time()

    # 创建临时目录
    os.makedirs(temp_dir, exist_ok=True)

    # 读取映射字典
    print("Reading reference dictionary ... ")
    referr_dict = pd.read_csv('large_referrer_dict.csv')
    id_to_value = dict(zip(referr_dict["referrer_id"], referr_dict["referrer_value"]))

    # 获取CPU核心数
    num_processes = cpu_count()
    print(f"Using {num_processes} processes for parallel processing")

    # 分块读取数据并行处理
    print("Processing data chunks ... ")
    chunk_size = 10000000
    processed_chunks = []

    with Pool(processes=num_processes) as pool:
        results = []
        chunk_num = 0

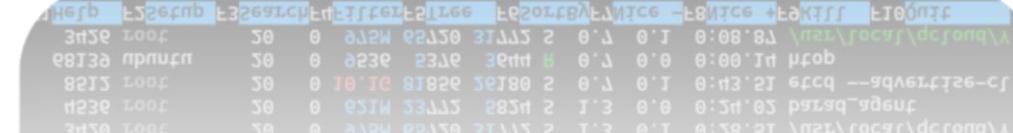
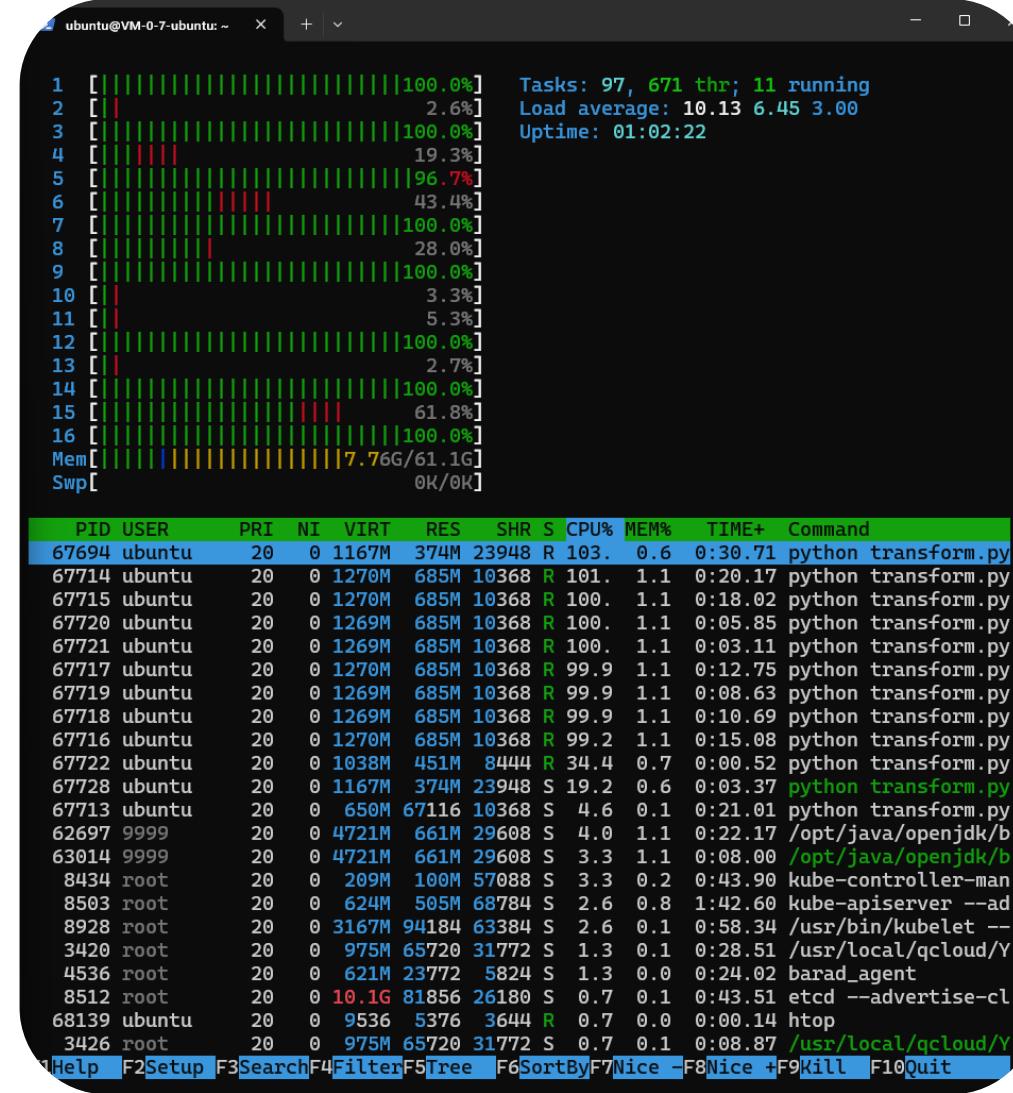
        # 读取数据并提交处理任务
        for chunk in pd.read_csv(input_file, sep=",", header=None,
                                 names=["web1_id", "web2_id", "similarity"],
                                 chunksize=chunk_size):
            results.append(pool.apply_async(process_chunk,
                                             args=(chunk.copy(), id_to_value, chunk_num)))
            chunk_num += 1

        # 收集处理结果
        for result in results:
            processed_chunks.append(result.get())

    # 合并所有临时文件
    print("Merging temporary files ... ")
    first_chunk = True

    while len(processed_chunks) > 0:
        if first_chunk:
            output_file = temp_dir + "/temp_file.csv"
            first_chunk = False
        else:
            output_file = temp_dir + "/temp_file.csv"

        with open(output_file, "a") as f:
            for chunk in processed_chunks:
                chunk.to_csv(f, index=False)
```



结果处理——显示最终结果

```
ubuntu@VM-0-7-ubuntu:~/all_tasks$ python transform.py
Reading reference dictionary...
Using 16 processes for parallel processing
Processing data chunks...
Merging temporary files...
Success! Processed 383157187 rows in 192.00s
```

```
20CC6221 6LOC62260 283T2AT8A IOM2 TU TAJ'002
```

```
ubuntu@VM-0-7-ubuntu:~/all_tasks$ wc -l large_result_final.csv
383157188 large_result_final.csv
ubuntu@VM-0-7-ubuntu:~/all_tasks$ head -n 20 large_result_final.csv
web1,web2,similarity
106895823676320711807,113290547889385032958,0.0252707581227436
10081612426229045070,109505227197640589104,0.0121155638397017
116337975768456202828,117344786984788205459,0.0611620795107033
113453874689606594546,116375847890292587830,0.0298755186721991
101320235608031237906,116385489466404408987,0.00836820083682
101320235608031237906,109974950487342042428,0.01666666666666666
101320235608031237906,115465838058347571798,0.01010101010101
115597296300763962023,118278972816201865535,0.03030303030303
100578645684682196743,101570646809005971541,0.0348623853211009
101320235608031237906,117713069287157269096,0.0108695652173913
101320235608031237906,114553528288744390401,0.0714285714285714
101320235608031237906,110207953520399557867,0.0454545454545454
101320235608031237906,115464467737653978226,0.0018315018315018
101136313452022372834,101214841967281809081,0.0445168295331161
100557819052469826667,108087777677620157134,0.1247947454844006
104243053732122719347,107974298339867918005,0.0580645161290322
110399178354646507311,118212287532576682860,0.0353356890459364
106895823676320711807,107169475135045361147,0.0151057401812688
106895823676320711807,115804671673356875627,0.0056818181818181
```

```
+ o|m o| + + ( )) + - + $ |  
J0E8a28533e3e350a7J80d,J280Hea7e332e8a2e5d,0.002e8J8J8T8J8T8  
J0E8a28533e3e350a7J80d,T07TeaHd2T320T23eJ7Hd,0.0J2T02aT0T8T5e88  
J103aa81832pHee20T3T,J185T538A2352Ae8e838e0,0.032332e8a0T2a23e8t  
J0T5H3023A35T52aT3Hd,T07a6aH2a833e8e6aT8002,0.0280eH2T2T2a0355  
J002p218T023Hpa833e8e0,T0808A1Ae2A2T0T2T3P,0.T5H2aT2p218p00e00
```

```
ubuntu@VM-0-7-ubuntu:~/all_tasks$ ls -l
```

```
total 35038084
-rw-rw-r-- 1 ubuntu ubuntu 2925518 May 16 21:28 large_referree_dict.csv
-rw-rw-r-- 1 ubuntu ubuntu 1966892 May 16 21:28 large_referrer_dict.csv
-rw-rw-r-- 1 ubuntu ubuntu 23957013430 May 19 16:24 large_result_final.csv
-rw-rw-r-- 1 ubuntu ubuntu 11917070025 May 19 16:12 task
-rw-rw-r-- 1 ubuntu ubuntu 3385 May 19 16:21 transform.py
```

```
+ o|m o| + + ( )) + - + $ |  
-LM-LM-TI SP078520/11807,115804672825556879827,0.005818181818181
```

```
ubuntu@VM-0-7-ubuntu:~/all_tasks$ tail -n 20 large_result_final.csv
100351380782421508832,112443872436342469463,0.0133333333333333
100351380782421508832,107486243467282501280,0.0010649627263045
107028565017991330517,112082748565101045822,0.009009009009009
111366756187576712078,11451664924445618945,0.0011286681715575
112889070874420371745,113832282534674162327,0.0045662100456621
111409417452299920769,116877582161876571800,0.0031645569620253
103506800973002590976,105411182022739904120,0.001081081081081
108420082152035649888,115875056630236947529,0.0158730158730158
100093468880409449436,102050273056135520452,0.0022123893805309
101875839400760158712,105731699847174230347,0.0032051282051282
101875839400760158712,113930216320253087485,0.0294117647058823
101875839400760158712,103627721981378204376,0.0011792452830188
101875839400760158712,111206226075949877358,0.0034013605442176
101875839400760158712,104359626267906791601,0.0163934426229508
101875839400760158712,109926473783208635057,0.03125
101875839400760158712,102718274791889610673,0.0294117647058823
101875839400760158712,110982884491550289466,0.0188679245283018
101875839400760158712,112312147769720349103,0.0083333333333333
103226166584311556931,109032868814925704218,0.00093896713615023
114122567270962169719,115011646412956679115,0.0033003300330033
```

```
+ o|m o| + + ( )) + - + $ |  
J1T5T52eA7a0e5Tea7a,TJ20T7eHd7J2a2e9a7T2,0.0033003300330033
J0352eTee2843TJ22e3T,10a0358e88JH2a210H5T8,0.000a38a6aT3eT2053
J0T81283aH001a0T28aT5,TJ53T5JH2a2e2a503H03,0.0083333333333333
J0T81283aH001a0T28aT5,TJ0a8388H2aT22038aH02,0.01886792583078
J0T81283aH001a0T28aT5,T03aT8a7a1a2T88a8T00a3,0.03a7T10p1028853
```

3. 问题处理

问题处理

```
-c
kubernetes-taskmanager.sh -Djobmanager.memory.jvm-overhead.min='429496736b' -Dpipeline.classpaths='' -Dtask
manager.resource-id='largejob-taskmanager-1-2' -Djobmanager.memory.off-heap.size='134217728b' -Dexecution.target=
'embedded' -Dweb.tmpdir='/tmp/flink-web-36c90164-d456-4939-8b14-0fc31c216414' -Djobmanager.rpc.port='6123' -Dpipe
line.jars='file:/opt/flink/largejob/inputfromfiledemo-1.0.jar' -Djobmanager.memory.jvm-metaspacesize='268435456b
'Djobmanager.memory.heap.size='3462817376b' -Djobmanager.memory.jvm-overhead.max='429496736b' -Dtaskmanager.me
mory.network.min='1073741824b' -Dtaskmanager.cpu.cores='15.0' -Dtaskmanager.memory.task.off-heap.size='0b' -Dtaskman
ager.memory.jvm-overhead.size='268435456b' -Dexternal-resources='none' -Dtaskmanager.memory.jvm-overhead.min='10737
41824b' -Dtaskmanager.memory.framework.off-heap.size='134217728b' -Dtaskmanager.memory.network.max='1073741824b
'Dtaskmanager.memory.framework.heap.size='134217728b' -Dtaskmanager.memory.managed.size='23514946296b' -Dtaskmanager.
memory.task.heap.size='33930241288b' -Dtaskmanager.numberoftaskSlots='15' -Dtaskmanager.memory.jvm-overhead.max='107
3741824b
Limits:
  cpu: 15
  memory: 56Gi
Requests:
  cpu: 15
  memory: 56Gi
Environment:
  _POD_NODE_ID: (v1:spec.nodeName)
  FLINK_TM_JVM_MEM_OPTS: -Xmx34064459016 -Xms34064459016 -XX:MaxDirectMemorySize=1207959552 -XX:MaxMetaspac
Size=268435456
Mounts:
  /opt/flink/conf from flink-config-volume (rw)
  /opt/flink/largejob from my-jar (rw)
  /var/run/secrets/kubernetes.io/serviceaccount from flink-token-wpqjx (ro)
Conditions:
  Type        Status
  PodScheduled  False
Volumes:
  my-jar:
    Type:      HostPath (bare host directory volume)
    Path:       /home/ubuntu/tableapp/upload
    HostPathType: Directory
  flink-config-volume:
    Type:      ConfigMap (a volume populated by a ConfigMap)
    Name:      flink-config-largejob
    Optional:  false
  flink-token-wpqjx:
    Type:      Secret (a volume populated by a Secret)
    SecretName: flink-token-wpqjx
    Optional:  false
QoS Class:  Guaranteed
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute for 300s
           node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type      Reason     Age           From           Message
  Warning   FailedScheduling  78s (x4 over 4m7s)  default-scheduler  0/3 nodes are available: 2 Insufficient memor
y, 3 Insufficient cpu.
ubuntu@VM-0-2-ubuntu:~/tableapp$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
flink-kubernetes-operator-8556b997b-dwvt  2/2     Running   0          79m
largejob-bdb87ddb-kfbk  1/1     Running   0          8m41s
largejob-taskmanager-1-1  1/1     Running   0          8m33s
largejob-taskmanager-1-2  0/1     Pending   0          8m33s
largejob-taskmanager-1-3  0/1     Pending   0          8m23s
largejob-taskmanager-1-4  0/1     Pending   0          8m23s
ubuntu@VM-0-2-ubuntu:~/tableapp$ |
```

跑flink作业时发现空闲节点
TaskManager节点只有一个

原因：

k8s对于JobManager和TaskManager随机分配
该节点被JobManager调度占用，不进
行计算

Master节点被自动打上污点，TM不会
分配到该节点；或master节点资源不
足，TM的pod一直处于pending状态

解决办法

```
taskManager:
  resource:
    memory: "51200m"
    cpu: 12 # 每个taskmanager分配12个CPU
  replicas: 3 # 启动2个taskmanager pod, 调度到2个节点
  podTemplate:
    spec:
      nodeSelector:
        flink-role: taskmanager
      tolerations:
        - key: "node-role.kubernetes.io/master"
          operator: "Exists"
          effect: "NoSchedule"
      containers:
        - name: flink-main-container
          volumeMounts:
            - name: my-jar
              mountPath: /opt/flink/largejob
      volumes:
        - name: my-jar
          hostPath:
            path: /home/ubuntu/tableapp/upload # 本地JAR所在目录
            type: Directory
  job:
    jarURI: local:///opt/flink/largejob/inputfromfiledemo-1.0.jar
    entryClass: cn.edu.shu.large_result # 填写作业的主类
    parallelism: 36
    b95ff6f2w: 30
```

移除master节点污点限制

```
kubectl taint node vm-0-9-ubuntu node-role.kubernetes.io/master:NoSchedule-
```

若想强制选择pod运行的节点

可以在yaml文件中添加节点选择标签，
只有当节点标签为taskmanager时才分配为TM

为节点打上标签

```
kubectl label nodes vm-0-14-ubuntu flink-role=taskmanager
```

显示节点标签

```
kubectl get nodes --show-labels
```

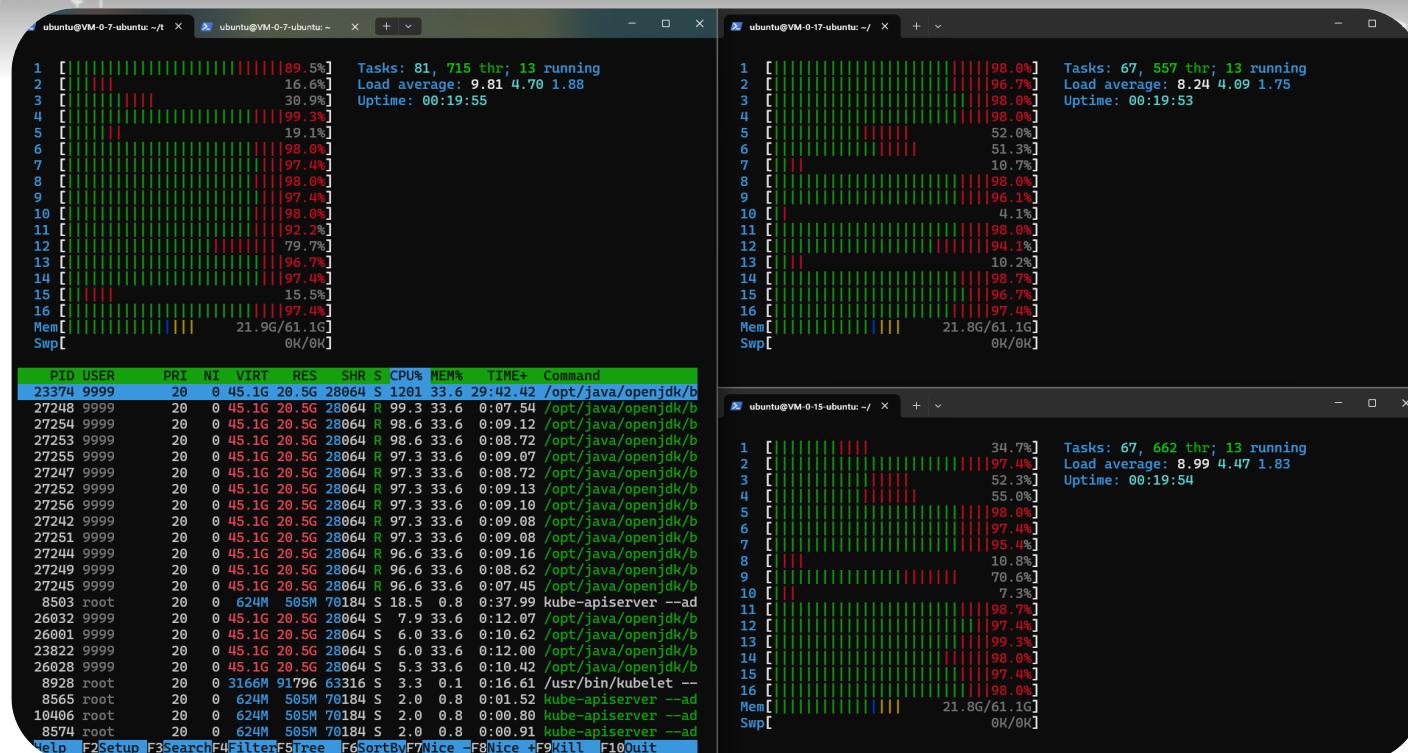
```
ubuntu@VM-0-7-ubuntu:~/tableapp$ kubectl get nodes --show-labels
NAME           STATUS  ROLES   AGE    VERSION  LABELS
vm-0-15-ubuntu Ready   <none>  15m   v1.18.0  beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,flink-role=taskmanager,kubernetes.io/arch=amd64,kubernetes.io/
hostname=vm-0-15-ubuntu,kubernetes.io/os=linux
vm-0-17-ubuntu Ready   <none>  15m   v1.18.0  beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,flink-role=taskmanager,kubernetes.io/arch=amd64,kubernetes.io/
hostname=vm-0-17-ubuntu,kubernetes.io/os=linux
vm-0-7-ubuntu  Ready   master   17m   v1.18.0  beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,flink-role=taskmanager,kubernetes.io/arch=amd64,kubernetes.io/
hostname=vm-0-7-ubuntu,kubernetes.io/os=linux,node-role.kubernetes.io/master=
ubuntu@VM-0-7-ubuntu:~/tableapp$
```

让TaskManager跑在3个节点上

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
flink-kubernetes-operator-8556b997b-mbbxg	2/2	Running	0	11m	10.244.2.4	vm-0-15-ubuntu	<none>	<none>
largejob-54b776cc5-9gshd	1/1	Running	0	5m6s	10.244.1.6	vm-0-17-ubuntu	<none>	<none>
largejob-taskmanager-1-1	1/1	Running	0	4m5s	10.244.2.5	vm-0-15-ubuntu	<none>	<none>
largejob-taskmanager-1-2	1/1	Running	0	4m5s	10.244.1.7	vm-0-17-ubuntu	<none>	<none>
largejob-taskmanager-1-3	1/1	Running	0	4m5s	10.244.0.4	vm-0-7-ubuntu	<none>	<none>
largejob-taskmanager-1-4	0/1	Pending	0	3m55s	<none>	<none>	<none>	<none>
largejob-taskmanager-1-5	0/1	Pending	0	3m55s	<none>	<none>	<none>	<none>
largejob-taskmanager-1-6	0/1	Pending	0	3m55s	<none>	<none>	<none>	<none>

```
ubuntu@VM-0-7-ubuntu:~/tableapp$
```

```
java -jar target/tableapp-0.1-1.jar
```



THANK YOU!

汇报人：22121630 汪江豪
22120721 阮金桐